

操作系统期末大作业

学 院：智能工程学院

姓 名：刘德鹏

学 号：18364059

邮 箱：liudp5@mail2.sysu.edu.cn

提交日期：2021/01/24

Xv6 lab: Multithreading/Uthread: switching between threads

仿照上下文切换的方式完成进程切换

在 trampoline.S 中的形式为

首先修改 uthread.c 中的 thread 定义如下

```
struct thread {
    /** My Implementation */
    uint64    ra;
    uint64    sp;
    // callee registers
    /** thread_switch needs to save/restore only the callee-save register
s. */
    uint64    s0;
    uint64    s1;
    uint64    s2;
    uint64    s3;
    uint64    s4;
    uint64    s5;
    uint64    s6;
    uint64    s7;
    uint64    s8;
    uint64    s9;
    uint64    s10;
    uint64    s11;

    char      stack[STACK_SIZE]; /* the thread's stack */
    int       state;              /* FREE, RUNNING, RUNNABLE */
};
```

模仿 proc.h 中进程上下文的切换，在 uthread_Switch.S 中补充线程上下文的切换：
修改 uthread_Switch.S 为

```
.text

/* Switch from current_thread to next thread_thread, and make
 * next_thread the current_thread.  Use t0 as a temporary register,
 * which should be caller saved. */

.globl thread_switch
thread_switch:
    /* YOUR CODE HERE */
    sd ra, 0(a0)
    sd sp, 8(a0)
    sd s0, 16(a0)
    sd s1, 24(a0)
    sd s2, 32(a0)
    sd s3, 40(a0)
    sd s4, 48(a0)
    sd s5, 56(a0)
    sd s6, 64(a0)
    sd s7, 72(a0)
    sd s8, 80(a0)
    sd s9, 88(a0)
    sd s10, 96(a0)
    sd s11, 104(a0)

    ld ra, 0(a1)
    ld sp, 8(a1)
    ld s0, 16(a1)
    ld s1, 24(a1)
    ld s2, 32(a1)
    ld s3, 40(a1)
    ld s4, 48(a1)
    ld s5, 56(a1)
    ld s6, 64(a1)
    ld s7, 72(a1)
    ld s8, 80(a1)
    ld s9, 88(a1)
    ld s10, 96(a1)
    ld s11, 104(a1)
    ret    /* return to ra */
```

修改 `thread_create`，使之能够记录线程的返回地址 `ra` 与栈地址 `sp`。其中，返回地址意味着当切换线程时，线程应该返回到什么哪个地址。这里应该是传入的函数的入口 `func`

```
void
thread_create(void (*func)())
{
    struct thread *t;

    for (t = all_thread; t < all_thread + MAX_THREAD; t++) {
        if (t->state == FREE) break;
    }
    t->state = RUNNABLE;
    // YOUR CODE HERE
    t->ra = (uint64)func;
    t->sp = (uint64)(t->stack + STACK_SIZE);
}
```

最后，在 `thread_schedule` 中添加 `thread_switch` 调用：

```
void
thread_schedule(void)
{
    struct thread *t, *next_thread;

    /* Find another runnable thread. */
    next_thread = 0;
    t = current_thread + 1;
    for(int i = 0; i < MAX_THREAD; i++){
        if(t >= all_thread + MAX_THREAD)
            t = all_thread;
        if(t->state == RUNNABLE) {
            next_thread = t;
            break;
        }
        t = t + 1;
    }
    if (next_thread == 0) {
        printf("thread_schedule: no runnable threads\n");
        exit(-1);
    }
    if (current_thread != next_thread) {          /* switch threads? */
        next_thread->state = RUNNING;
        t = current_thread;
        current_thread = next_thread;
        /* YOUR CODE HERE
```

```

    * Invoke thread_switch to switch from t to next_thread:
    * thread_switch(??, ??);
    */
    thread_switch((uint64)t, (uint64)next_thread);
} else
    next_thread = 0;
}

```

测试结果

```

liudp5@ubuntu: ~/Desktop/xv6-riscv-fall19
xv6 kernel is booting
virtio disk init 0
hart 1 starting
hart 2 starting
init: starting sh
$ uthread
thread_a started
thread_b started
thread_c started
thread_c 0
thread_a 0
thread_b 0
thread_c 1
thread_a 1
thread_b 1
thread_c 2
thread_a 2
thread_b 2
thread_c 3
thread_a 3
thread_b 3

```

```

thread_c 94
thread_a 94
thread_b 94
thread_c 95
thread_a 95
thread_b 95
thread_c 96
thread_a 96
thread_b 96
thread_c 97
thread_a 97
thread_b 97
thread_c 98
thread_a 98
thread_b 98
thread_c 99
thread_a 99
thread_b 99
thread_c: exit after 100
thread_a: exit after 100
thread_b: exit after 100
thread_schedule: no runnable threads
$

```

Xv6 lab: Lock/Memory allocator

由于有多个 CPU
实现一个多 free list 的初始化

```
struct kmem{
    struct spinlock lock;
    struct run *freelist;
} kmem;

struct kmem kmems[NCPU]; //多个 kmem
void
kinit() //实现多 free list 的初始化
{
    push_off();
    int currentid = cpuid();
    pop_off();

    printf("# cpuId:%d \n", currentid);
    /* 初始化 NCPU 个锁 */
    for (int i = 0; i < NCPU; i++)
    {
        initlock(&kmems[i].lock, "kmem");
    }
    //initlock(&kmem.lock, "kmem");
    freerange(end, (void*)PHYSTOP);
}
```

实现插入和弹出，将 free list 的插入 push() 和弹出 pop() 操作封装为函数

```
struct run* trypopr(int id){
    struct run *r;
    r = kmems[id].freelist;
    if(r)
        kmems[id].freelist = r->next;
    return r;
}

void trypushr(int id, struct run* r){
    if(r){
        r->next = kmems[id].freelist;
        kmems[id].freelist = r;
    }
    else
    {
        panic("cannot push null run");
    }
}
```

修改 kfree()函数, 查看空余 CPU, 用 cpuid()获取相应的 id, 然后将被释放的块插入相应的 free list

```
void
kfree(void *pa)
{
    struct run *r;

    if(((uint64)pa % PGSIZE) != 0 || (char*)pa < end || (uint64)pa >= PHYSTOP)
        panic("kfree");
    // Fill with junk to catch dangling refs.
    memset(pa, 1, PGSIZE);
    r = (struct run*)pa;

    push_off();
    int currentid = cpuid();
    pop_off();

    acquire(&kmems[currentid].lock);
    trypushr(currentid, r);
    release(&kmems[currentid].lock);
}
```

当一个 CPU 的 free list 为空, 但是另一个 CPU 的 free list 还有空闲块时, 我们就应该从有空闲 free list 的 CPU 处“偷”一个空闲内存块。

在 kalloc()函数中, 首先检查当前 CPU 的 free list, 如果当前 CPU 的 free list 不为空, 那就直接 pop 一个 run 出来, 将之初始化后, 返回给调用者; 如果当前 CPU 对应的 free list 为空, 就需要去查询其它 CPU 对应的 free list, 找到空闲的块 r 后, 按照如下步骤操作:

- 1.将 r 插入自己的 free list 中;
- 2.将 r 从自己的 free list 中弹出, 将之初始化, 返回给调用者;

```
void *
kalloc(void)
{
    struct run *r;
    int issteal = 0;/** 标识是否为偷盗 */
    push_off();
    int currentid = cpuid();
    pop_off();

    acquire(&kmems[currentid].lock);

    r = trypopr(currentid);
    /**
     * 将 id 的一块 free page 卸下, 然后给 currentid
```

```

* 这个过程中经历了:
*
* 1.卸下 id 的 free page
* 2.为 current id 的 freelist 添加该 page
* 3.将 current id 的 freelist 中的该 page 卸载掉
* 4.返回该 page
*
* 整个过程完成了 current id 偷盗 id 的 free page 的行为
*/
if(!r){
    //printf("oops out of memory\n");
    for (int id = 0; id < NCPU; id++){
        /* steal first run */
        if(id != currentid){
            /** 锁住 id 的 freelist, 此时不让其他 cpu 访问 */
            if(kmems[id].freelist){
                acquire(&kmems[id].lock);
                /** 卸下 id 的 free page */
                r = trypopr(id);
                /** 为 currentid 的 freelist 添加一个 run */
                trypushr(currentid, r);
                issteal = 1;
                release(&kmems[id].lock);
                break;
            }
        }
        //printf("\n");
    }
}
/** 如果是偷盗的, 则把 currentid 的 freelist 释放出来 */
if(issteal)
    r = trypopr(currentid);

release(&kmems[currentid].lock);

if(r){
    //printf("currentid: %d, r: %p\n", currentid, r);
    memset((char*)r, 5, PGSIZE); // fill with junk
}
/** 返回该 page */
//printf("issteal: %d \n", issteal);
return (void*)r;
}

```

测试结果

```
liudp5@ubuntu: ~/Desktop/xv6-riscv-fall19
# cpuId:0
virtio disk init 0
hart 2 starting
hart 1 starting
init: starting sh
$ kalloc_test
start test0
test0 results:
=== lock kmem/bcache stats
lock: kmem: #test-and-set 0 #acquire() 179117
lock: kmem: #test-and-set 0 #acquire() 200027
lock: kmem: #test-and-set 0 #acquire() 20892
=== top 5 contended locks:
lock: cons: #test-and-set 0 #acquire() 15
lock: cons: #test-and-set 0 #acquire() 15
lock: cons: #test-and-set 0 #acquire() 15
lock: cons: #test-and-set 0 #acquire() 15
lock: cons: #test-and-set 0 #acquire() 15
test0 OK
start test1
total allocated number of pages: 200000 (out of 32768)
test1 OK
$
```

Xv6 lab: Lock/Buffer cache

首先修改 bcache，以使其支持哈希桶结构：

```
struct {
    struct spinlock lock;
    struct buf buf[NBUF];           // block[30]
    /** 实则循环双向链表 */
    // Linked list of all buffers, through prev/next.
    // head.next is most recently used.
    //struct buf head;

    struct buf buckets[NBUKETS];
    struct spinlock bucketslock[NBUKETS];
} bcache;
```

初始化 bcache:

```
void
binit(void)
{
    struct buf *b;
    /** 在 head 头插入 b */
    initlock(&bcache.lock, "bcache");
    for (int i = 0; i < NBUKETS; i++)
    {
```



```

    initlock(&bcache.bucketslock[i], "bcache.bucket");
    bcache.buckets[i].prev = &bcache.buckets[i];
    bcache.buckets[i].next = &bcache.buckets[i];
}

for (b = bcache.buf; b < bcache.buf + NBUF; b++)
{
    int hash = getHb(b);
    b->time_stamp = ticks;
    b->next = bcache.buckets[hash].next;
    b->prev = &bcache.buckets[hash];
    initsleeplock(&b->lock, "buffer");
    bcache.buckets[hash].next->prev = b;
    bcache.buckets[hash].next = b;
}
}

```

bget(), 其实完成这一部分完全就是将原代码改写为支持哈希桶即可。需要注意的是：在 bget 函数里，如果找到相应缓冲区的话，就简单返回就行，但如果没找到，就去其他 hash 桶里偷个来放自己所属的缓冲区里，如果别的 hash 桶里没有的话就报错

```

static struct buf*
bget(uint dev, uint blockno)
{
    int hash = getH(blockno);
    struct buf *b;
    /**
     *
     * My modification
     */
    acquire(&bcache.bucketslock[hash]);

    for(b = bcache.buckets[hash].next; b != &bcache.buckets[hash]; b = b->next){
        if(b->dev == dev && b->blockno == blockno){
            b->time_stamp = ticks;
            b->refcnt++;
            //printf("## end has \n");
            release(&bcache.bucketslock[hash]);
            acquiresleep(&b->lock);
            return b;
        }
    }
}

```

```

    // If there is no cached buffer for the given sector, bget must make
    one, possibly reusing a buffer
    // that held a different sector. It scans the buffer list a second ti
    me, looking for a buffer that
    // is not in use
    // (b->refcnt = 0); any such buffer can be used. Bget edits the buffe
    r metadata to record the new
    // device and sector number and acquires its sleep-lock. Note that th
    e assignment b->valid = 0
    // ensures that bread will read the block data from disk rather than
    incorrectly using the buffer's
    // previous contents.
    // Not cached; recycle an unused buffer.

/**
 *
 * My modification
 */
for (int i = 0; i < NBUKETS; i++)
{
    if(i != hash){
        acquire(&bcache.bucketslock[i]);
        for(b = bcache.buckets[i].prev; b != &bcache.buckets[i]; b = b->p
rev){
            if(b->refcnt == 0){
                b->time_stamp = ticks;
                b->dev = dev;
                b->blockno = blockno;
                b->valid = 0;    //important
                b->refcnt = 1;

                /** 将 b 脱出 */
                b->next->prev = b->prev;
                b->prev->next = b->next;

                /** 将 b 接入 */
                b->next = bcache.buckets[hash].next;
                b->prev = &bcache.buckets[hash];
                bcache.buckets[hash].next->prev = b;
                bcache.buckets[hash].next = b;
                //printf("## end alloc: hash: %d, has: %d\n", hash,i);
                release(&bcache.bucketslock[i]);
                release(&bcache.bucketslock[hash]);
                acquiresleep(&b->lock);
            }
        }
    }
}

```

```

        return b;
    }
}
release(&bcache.bucketslock[i]);
}
}
panic("bget: no buffers");
}

```

实现 `brelse()`，不需要采用任何 `lock`，利用时间戳机制就能解决问题

```

void
brelse(struct buf *b)
{
    //printf("#----- brelse! -----
    -----\n");

    if(!holdingsleep(&b->lock))
        panic("brelse");

    releasesleep(&b->lock);
    int blockno = getHb(b);
    b->time_stamp = ticks;
    if(b->time_stamp == ticks){
        b->refcnt--;
        if(b->refcnt == 0){
            /** 将 b 脱出 */
            b->next->prev = b->prev;
            b->prev->next = b->next;

            /** 将 b 接入 */
            b->next = bcache.buckets[blockno].next;
            b->prev = &bcache.buckets[blockno];
            bcache.buckets[blockno].next->prev = b;
            bcache.buckets[blockno].next = b;
        }
    }
}

```

测试结果

```

hart 2 starting
hart 1 starting
init: starting sh
$ bcachetest
start test0
test0 results:
--- lock kmem/bcache stats
lock: kmem: #fetch-and-add 0 #acquire() 32982
lock: kmem: #fetch-and-add 0 #acquire() 59
lock: kmem: #fetch-and-add 0 #acquire() 39
lock: bcache: #fetch-and-add 0 #acquire() 2481
lock: bcache: #fetch-and-add 0 #acquire() 1448
lock: bcache: #fetch-and-add 0 #acquire() 2496
lock: bcache: #fetch-and-add 0 #acquire() 2170
lock: bcache: #fetch-and-add 0 #acquire() 3179
lock: bcache: #fetch-and-add 0 #acquire() 3176
lock: bcache: #fetch-and-add 0 #acquire() 3392
lock: bcache: #fetch-and-add 0 #acquire() 3219
lock: bcache: #fetch-and-add 0 #acquire() 4781
lock: bcache: #fetch-and-add 0 #acquire() 3601
lock: bcache: #fetch-and-add 0 #acquire() 2736
lock: bcache: #fetch-and-add 0 #acquire() 2475
lock: bcache: #fetch-and-add 0 #acquire() 1474
--- top 5 contended locks:
lock: virtio_disk: #fetch-and-add 191397 #acquire() 1378
lock: proc: #fetch-and-add 113049 #acquire() 87151
lock: proc: #fetch-and-add 56438 #acquire() 86747
lock: proc: #fetch-and-add 54043 #acquire() 86726
lock: proc: #fetch-and-add 48809 #acquire() 86726
tot= 0
test0: OK
start test1
test1 OK

```

Xv6 lab: File System/Large files

- 1.为 inode 添加一个二级索引，使之能够支持分配更大的文件；
- 2.学会如何创建 symbol link 软链接

为 inode 添加一个 Double Indirect 索引块，这样每一个 inode 就能够支持 $11 + 256 + 256 * 256 = 65803$ 个数据块的文件了。修改宏定义如下

```

/** direct blocks */
#define NDIRECT 11
/** indirect blocks */
#define SINGLEDIRECT (BSIZE / sizeof(uint))
#define NINDIRECT (SINGLEDIRECT + SINGLEDIRECT * BSIZE / sizeof(uint))
#define MAXFILE (NDIRECT + NINDIRECT)

```

其中，NDIRECT 代表直接索引指向块的数量；SINGLEDIRECT 表示一级索引能够指向的块的数量；NINDIRECT 索引代表一级和二级索引一共可指向的块的数量。

修改 kernel/param.h 中的 FSSIZE 宏定义为 200000

```
#define FSSIZE      200000  // size of file system in blocks
```

模仿 bmap 原有分配块的操作，实现二级索引列表

```
static uint
bmap(struct inode *ip, uint bn)
{
    /**
     * bmap() deals with two kinds of block numbers.
     * The bn argument is a "logical block number" --
     * a block number within the file, relative to the start of the file.
     *
     * The block numbers in ip->addrs[], and the argument to bread(),
     * are disk block numbers.
     * You can view bmap() as
     *
     * mapping a file's logical block numbers into disk block numbers.
     */
    // printf("-----\n");
    // printf("bn1: %d\n", bn);
    // NINDIRECT 256
    // printf("NINDIRECT: %d\n", NINDIRECT); //256
    uint addr, *a, *a2;
    struct buf *bp, *bp2;
    /** bn 是相对 file (inode) 的虚拟编号 */
    /** 直接返回块 ip->addrs[bn] */
    if(bn < NDIRECT){
        if((addr = ip->addrs[bn]) == 0)
            ip->addrs[bn] = addr = balloc(ip->dev);
        return addr;
    }
    /** 否则减去 NDIRECT */
    bn -= NDIRECT;
    //11 + 256 + 256 * 256
    if(bn < NINDIRECT){
        if(bn < SINGLEDIRECT){
            // Load indirect block, allocating if necessary.
            if((addr = ip->addrs[NDIRECT]) == 0)
                ip->addrs[NDIRECT] = addr = balloc(ip->dev);
            /** Return a locked buf with the contents of the indicated block.
             */
            bp = bread(ip->dev, addr);
            /**
             * 为何 #define NINDIRECT (BSIZE / sizeof(uint)) ?
             */
        }
    }
}
```

* 因为一个 buf 有 BSIZE 个 data，一个 block 地址为 64 位，故刚好存储 256 个 block 指针，NINDIRECT = 256

```
*/
/** 一个文件描述符 */
a = (uint*)bp->data;
if((addr = a[bn]) == 0){
    a[bn] = addr = balloc(ip->dev);
    log_write(bp);
}
brelse(bp);
}
else
{
    /**
     * 实现 double-indirect
     * 结构:
     * addr -> 256 个 single-indirect -> 256 * 256 个 block
     * 256 + 11 = 267 0 ~ 266 为前面的, 267 开始后面为后面的
     * */
    bn -= SINGLEDIRECT;
    /** 宏在做乘除运算时记得加括号 55555555555555555555 */
    /** single-indirect 索引*/
    int single_indirect_index = bn / SINGLEDIRECT;
    /** single-indirect 内部相对索引 */
    int relative_offset_bn = bn % SINGLEDIRECT;
    /** 下标 12 是 double-indirect */
    int pos = NDIRECT + 1;
    if((addr = ip->addrs[pos]) == 0)
        ip->addrs[pos] = addr = balloc(ip->dev);

    bp = bread(ip->dev, addr);
    a = (uint*)bp->data;
    if((addr = a[single_indirect_index]) == 0){
        printf("bn: %d, addr: %p\n", bn, addr);
        a[single_indirect_index] = addr = balloc(ip->dev);
        log_write(bp);
    }
    brelse(bp);
    bp2 = bread(ip->dev, addr);
    a2 = (uint*)bp2->data;
    if((addr = a2[relative_offset_bn]) == 0){
        a2[relative_offset_bn] = addr = balloc(ip->dev);
        log_write(bp2);
    }
}
```

```

        brelse(bp2);
    }
    printf("-----\n");
    return addr;
}
panic("bmap: out of range");
}

```

将 itrunc 补充完整，以释放 inode 下的所有数据块，同样，做法参考原本的 itrunc 即可

```

static void
itrunc(struct inode *ip)
{
    int i, j;
    struct buf *bp, *bp2;
    uint *a, *a2;
    /** Free 掉所有 Direct */
    for(i = 0; i < NDIRECT; i++){
        if(ip->addrs[i]){
            bfree(ip->dev, ip->addrs[i]);
            ip->addrs[i] = 0;
        }
    }
    /** Free 掉 single-direct */
    if(ip->addrs[NDIRECT]){
        bp = bread(ip->dev, ip->addrs[NDIRECT]);
        a = (uint*)bp->data;
        /**
         * 修改 NINDIRECT 为 SINGLEDIRECT
         * for(j = 0; j < NINDIRECT; j++){
         *     if(a[j])
         *         bfree(ip->dev, a[j]);
         * }
         */
        for (j = 0; j < SINGLEDIRECT; j++)
        {
            /** code */
            if(a[j])
                bfree(ip->dev, a[j]);
        }

        brelse(bp);
        bfree(ip->dev, ip->addrs[NDIRECT]);
        ip->addrs[NDIRECT] = 0;
    }
}

```

```

/** Free 掉 double-direct */
if(ip->addrs[NDIRECT + 1]){
    printf("free double\n");
    int pos = NDIRECT + 1;
    bp = bread(ip->dev, ip->addrs[pos]);
    a = (uint*)bp->data;
    int number_of_single_direct = BSIZE / sizeof(uint);
    for (i = 0; i < number_of_single_direct; i++)
    {
        if(a[i]){
            bp2 = bread(ip->dev, a[i]);
            a2 = (uint *)bp2->data;
            for (j = 0; j < SINGLEDIRECT; j++)
            {
                if(a2[j])
                    bfree(ip->dev, a2[j]);
            }
            brelse(bp2);
            bfree(ip->dev, a[i]);
            a[i] = 0;
        }
    }
    brelse(bp);
    bfree(ip->dev, ip->addrs[pos]);
    ip->addrs[pos] = 0;
}
ip->size = 0;
iupdate(ip);
}

```

测试结果

```

$ bigfile
.....
.....
.....
.....
wrote 65803 blocks
bigfile done; ok

```